GENERAL ASSEMBLY

DATA SCIENCE
CLASS 1: INTRO TO DATA SCIENCE

0. WHAT IS A DATA SCIENTIST?
I. HOW DATA SCIENTISTS ADD VALUE
II. THE DATA MINING WORKFLOW
III. QUALITIES OF A GOOD DATA SCIENTIST

0. WHAT IS A DATA SCIENTIST?

WHAT IS YOUR DEFINITION?

WHAT IS A DATA SCIENTIST?

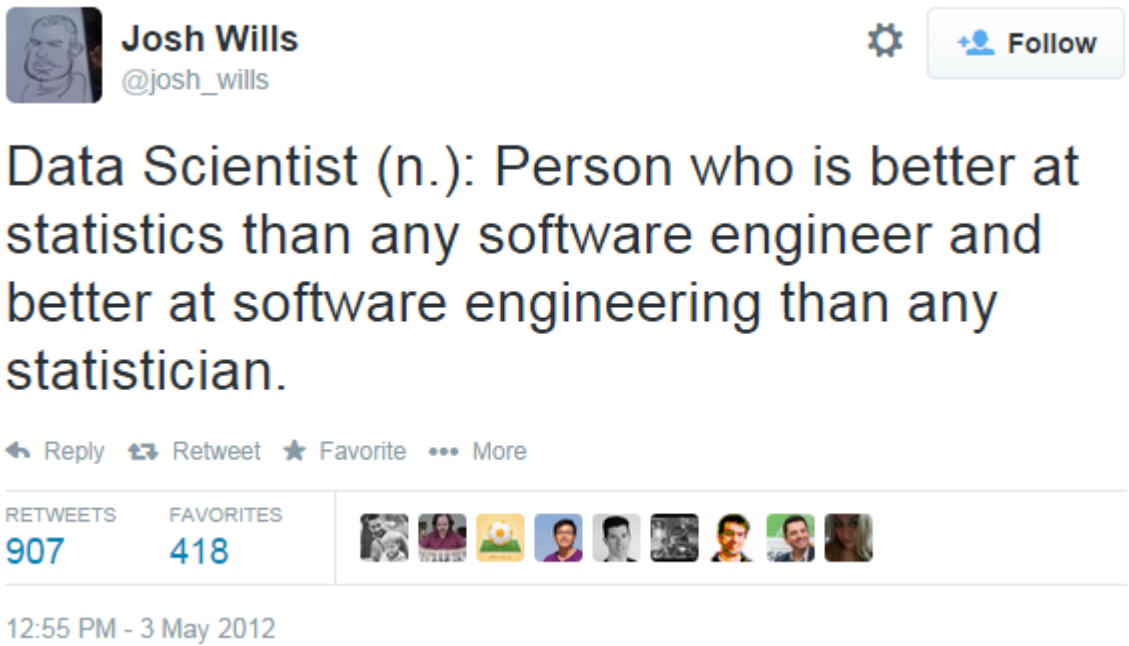
Josh Wills
@josh_wills
Follow
Data Scientist (n.): Person who is better at statistics than any software engineer and better at software engineering than any statistician.
Reply
Retweet
Favorite
More
RETWEETS
907
FAVORITES
418
12:55 PM - 3 May 2012

Data Scientists solve complex problems using data mining techniques

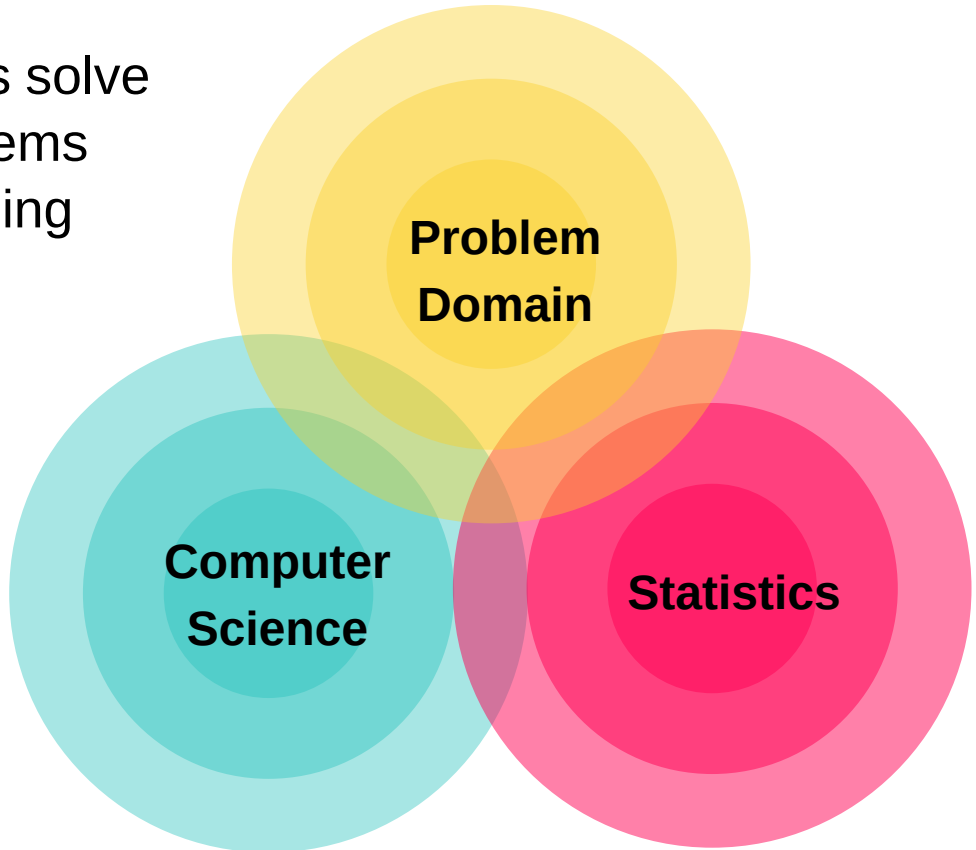

Data Scientists solve complex problems using data mining techniques

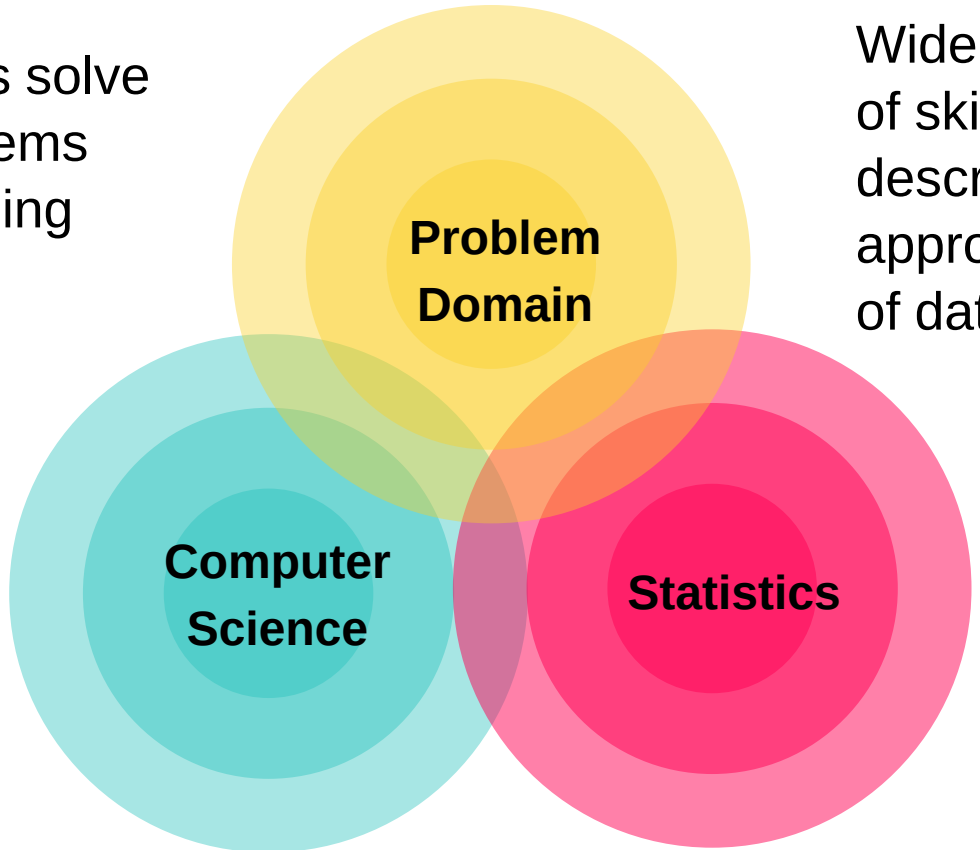

Wide variance in terms of skillsets: many job descriptions are more appropriate for a team of data scientists

I. HOW DATA SCIENTISTS ADD VALUE

5 Minutes:

Take 5 minutes, search the internet, and write down as many examples as you can of data scientists in action.

5 Minutes:

In a small group, introduce yourselves. As a group, decide how to best “cluster” your examples around how the data scientist adds value. You can have as many clusters as you want! Make sure you have a label for each cluster.

Data mining techniques generally add value by doing one of three things:

1) Predicting the bad
2) Identifying the good
3) Automating existing processes

Data scientists can be found within many fields: let's look at some additional examples to motivate this course.

Source: https://www.youtube.com/watch?v=fPzmnRj671Y

EXAMPLE #1: PREDICTING NEONATAL INFECTION

Problem: Children born prematurely are at high risk of developing infections, many of which are not detected until after the baby is sick

Goal: Detect subtle patterns in the data that predicts infection before it occurs

Data: 16 vital signs such as heart rate, respiration rate, blood pressure, etc…

Impact: Model is able to predict the onset of infection 24 hours before the traditional symptoms of infection appear

Image: http://www.babycaretips4u.com/wp-content/uploads/2014/03/premature-baby.jpg
Case Study: http://www.amazon.com/Big-Data-Revolution-Transform-Think/dp/0544002695

EXAMPLE #2: AUTOMATING GOVERNMENT PAPER-PUSHING

Problem: Processing disability claims at the Social Security Administration is a time-intensive process, with many claims taking over 2 years to adjudicate

Goal: Automate the approval of a subset of the "simplest" disability claims

Data: Free text in the claims form

Impact: Able to fully automate 20% of the simplest claims. Rating accuracy of the algorithm is higher than the average claims examiner.

Case Study: http://datamininglab.com/images/case-studies/ERI_Text_Mining_SSA_Claims_for_Disability_Approval.pdf

II. THE DATA MINING WORKFLOW

0. Define the problem / question
I. Identify and collect data
II. Explore and prepare data
III. Build and evaluate model
IV. Communicate results

0. DEFINE THE PROBLEM / QUESTION

Can I predict infection before it occurs?

Can I predict claim approval from the start of the process?

I. IDENTIFY AND COLLECT DATA

Vital Areas: Heart Rate, Blood Pressure, etc…

Want to collect all data on the claim form (mostly free text)

II. EXPLORE AND PREPARE DATA

Aggregate data at the minute level

Cluster like words

III. BUILD AND EVALUATE MODELS

Compare Decision Tree with Logistic Regression

Start with Naïve Bayes Classifier

IV. COMMUNICATE RESULTS

Create custom dashboard for doctors and nurses

Create report and dashboard proof of concept

III. QUALITIES OF A GOOD DATA SCIENTIST

ASKS RATIONAL QUESTIONS

Law of Small Numbers

The Anchor Effect

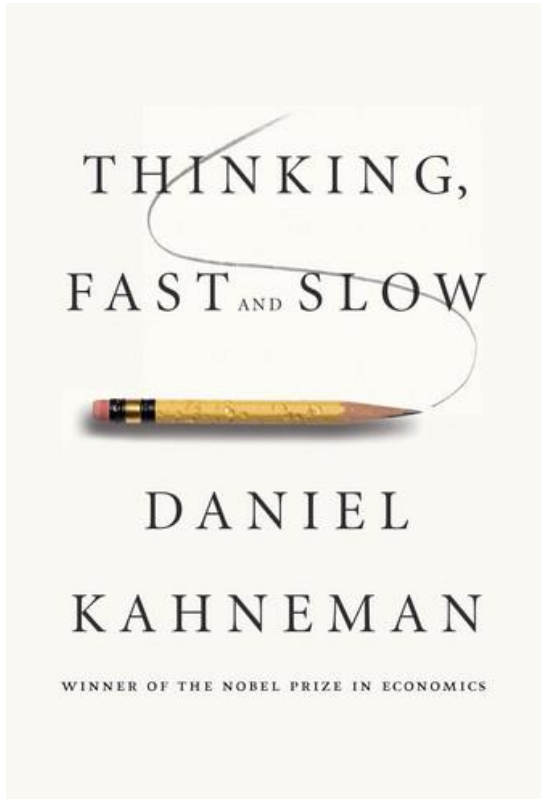

RETAINS INTELLECTUAL HUMILITY

COMMUNICATES CLEARLY

UNDERSTANDS THE PROS & CONS OF DIFFERENT TECHNIQUES

STATISTICIANS, LIKE ARTISTS, HAVE THE BAD HABIT OF FALLING IN LOVE WITH THEIR MODELS – GEORGE BOX